

▼ Used_Car_Marketing_Analysis

KNN_Regression,Linear Regression,Random Forest

Download Dataset

<https://www.kaggle.com/orgesleka/used-cars-database/discussion>

Import Libraries and Load Dataset

```
from google.colab import drive
drive.mount('/content/drive')
```

🔗 Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id

Enter your authorization code:

.....

Mounted at /content/drive

```
#Load libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Load dataset
df=pd.read_csv('../content/drive/My Drive/DSProject/Germany_Used_Car
```

🔗 /usr/local/lib/python3.6/dist-packages/IPython/core/interactiveshell.py:2718: D
interactivity=interactivity, compiler=compiler, result=result)

Exploratory Analysis data

```
# Dataset preview
df.head()
```

🔗

	dateCrawled	name	price	vehicleType
0	3/21/2016 21:46	Opel_Corsa_1.2_16V__2_HAND_KLIMA_8_FACH_T?V...	2238	small car
1	4/4/2016 23:48	Mercedes-Benz_E_220_T_CDI_Avantgarde	12500	station wagon
2	3/17/2016 0:46	BMW_325_xi_E92_Coupe__EZ_12/2006__6__Gang_Sch...	13299	coupe
3	3/29/2016 18:51	BMW_520d_Touring_Xenon_Navi+_PCD_Sport_Comf._1...	17200	station wagon
4	3/28/2016 16:45	BMW_635_CSI_Schaltgetriebe_original_Fahrzeug_2...	8000	coupe

```
# Review all column index
df.columns.values
```

```
↳ array(['dateCrawled', 'name', 'price', 'vehicleType',
        'yearOfRegistration', 'yearOfUsing', 'gearbox', 'powerPS', 'model',
        'kilometer', 'monthOfRegistration', 'fuelType', 'brand',
        'notRepairedDamage', 'dateCreated', 'postalCode', 'lastSeen'],
        dtype=object)
```

```
# Number of rows and columns
print(df.shape)
```

```
↳ (242109, 17)
```

```
# Data types
df.info()
```

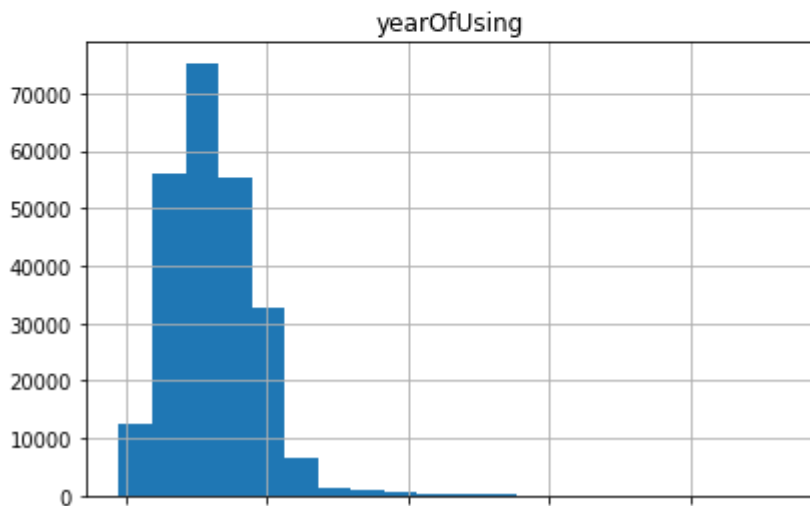

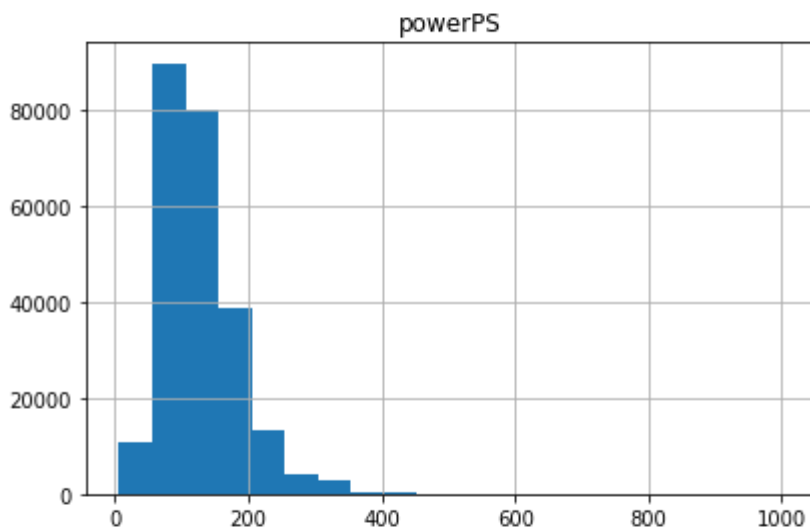

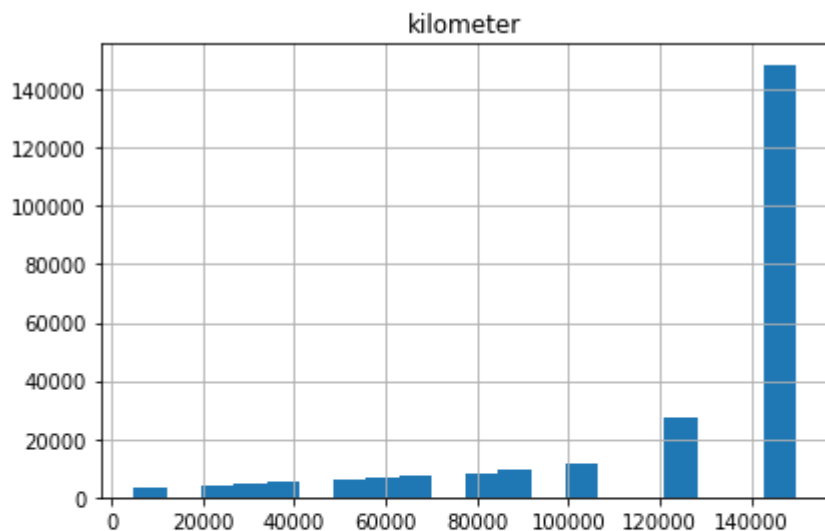
```
↳
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 242109 entries, 0 to 242108
Data columns (total 17 columns):
dateCrawled      242109 non-null object
name            242109 non-null object
price           242109 non-null object
vehicleType     242109 non-null object
yearOfRegistration 242109 non-null int64
yearOfUsing     242109 non-null int64
gearbox         242109 non-null object
powerPS        242109 non-null int64
model          242109 non-null object
kilometer       242109 non-null int64
monthOfRegistration 242109 non-null int64
fuelType       242109 non-null object
brand          242109 non-null object
notRepairedDamage 242109 non-null object
dateCreated     242109 non-null object
postalCode     242109 non-null object
lastSeen       242109 non-null object
dtypes: int64(5), object(12)
memory usage: 31.4+ MB
```

```
# Visualization for data statistic distribution
df.hist(figsize = (15,15),bins=20)
```



```
array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7f1108ed9710>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f1108eb1c50>],  
      [<matplotlib.axes._subplots.AxesSubplot object at 0x7f1108e71240>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f1108e217f0>],  
      [<matplotlib.axes._subplots.AxesSubplot object at 0x7f1108dd1da0>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f1108d8c390>]],  
      dtype=object)
```



0

20

40


60

80

```
# Looking at the frequency of most popular Car Brand
plt.rcParams['figure.figsize'] = (18, 7)
color = plt.cm.copper(np.linspace(0, 1, 40))
df.brand.value_counts().head(40).plot.bar(color = color)
plt.title('frequency of most popular used car', fontsize = 20)
plt.xticks(rotation = 90 )
plt.grid()
plt.show()
```



frequency of most popular us



▼ Feature Engineering

```
# Review all column index
df.columns.values
```



```
array(['dateCrawled', 'name', 'price', 'vehicleType',
      'yearOfRegistration', 'yearOfUsing', 'gearbox', 'powerPS', 'model',
      'kilometer', 'monthOfRegistration', 'fuelType', 'brand',
      'notRepairedDamage', 'dateCreated', 'postalCode', 'lastSeen'],
      dtype=object)
```

```
# Check null data
```

```
df.isnull().sum()
```

```
↳ dateCrawled      0
   name            0
   price           0
   vehicleType     0
   yearOfRegistration 0
   yearOfUsing     0
   gearbox         0
   powerPS         0
   model           0
   kilometer       0
   monthOfRegistration 0
   fuelType        0
   brand           0
   notRepairedDamage 0
   dateCreated     0
   postalCode      0
   lastSeen        0
   dtype: int64
```

```
# We found 'price' is mixed data types with string ('Lower Saxony')
```

```
# Only four rows has 'Lower Saxony' so that we drop them.
```

```
df_remove = df[df.price == 'Lower Saxony']
```

```
df= df.drop(df_remove.index)
```

```
# Convert price to integer
```

```
df['price']=df['price'].astype(int)
```

```
# Convert to str
```

```
columns = ['yearOfRegistration', 'monthOfRegistration', 'postalCode']
```

```
for x in columns:
```

```
    df[x]=df[x].astype(str)
```

```
# Data Statistics
```

```
df.describe()
```

```
↳
```

	price	yearOfUsing	powerPS	kilometer
count	242105.000000	242105.000000	242105.000000	242105.000000
mean	6752.196968	12.385717	129.297173	123789.264988
std	7656.033019	6.297437	61.339541	39741.321291
min	200.000000	-1.000000	6.000000	5000.000000
25%	1700.000000	8.000000	86.000000	100000.000000
50%	4000.000000	12.000000	116.000000	150000.000000
75%	8950.000000	16.000000	160.000000	150000.000000
max	99999.000000	93.000000	1000.000000	150000.000000


brand

```
df.brand.describe()
```

```
count      242105
unique         39
top    volkswagen
freq       51443
Name: brand, dtype: object
```

```
# Visualising all the brand
plt.figure(figsize=(15,8))
ax = sns.countplot(df.brand)
ax.set_xticklabels(ax.get_xticklabels(), rotation=90, ha="right")
plt.show()
```

```
↳
```



```
y = df.brand.value_counts().to_frame()  
y.head(10)
```



	brand
volkswagen	51443
bmw	29175
mercedes_benz	25526
opel	24321
audi	22699
ford	16094
renault	10546
peugeot	7412
fiat	5093
seat	4758

Visualization for Brand

read new data frame brand1, will show top 15 best sale used car
 and1=df.brand.value_counts().head(15).reset_index().rename(columns={

```
brand1['Car'] = 'CAR'
Car = brand1.truncate(before = -1, after = 15)
```

```
Car
import networkx as nx
```

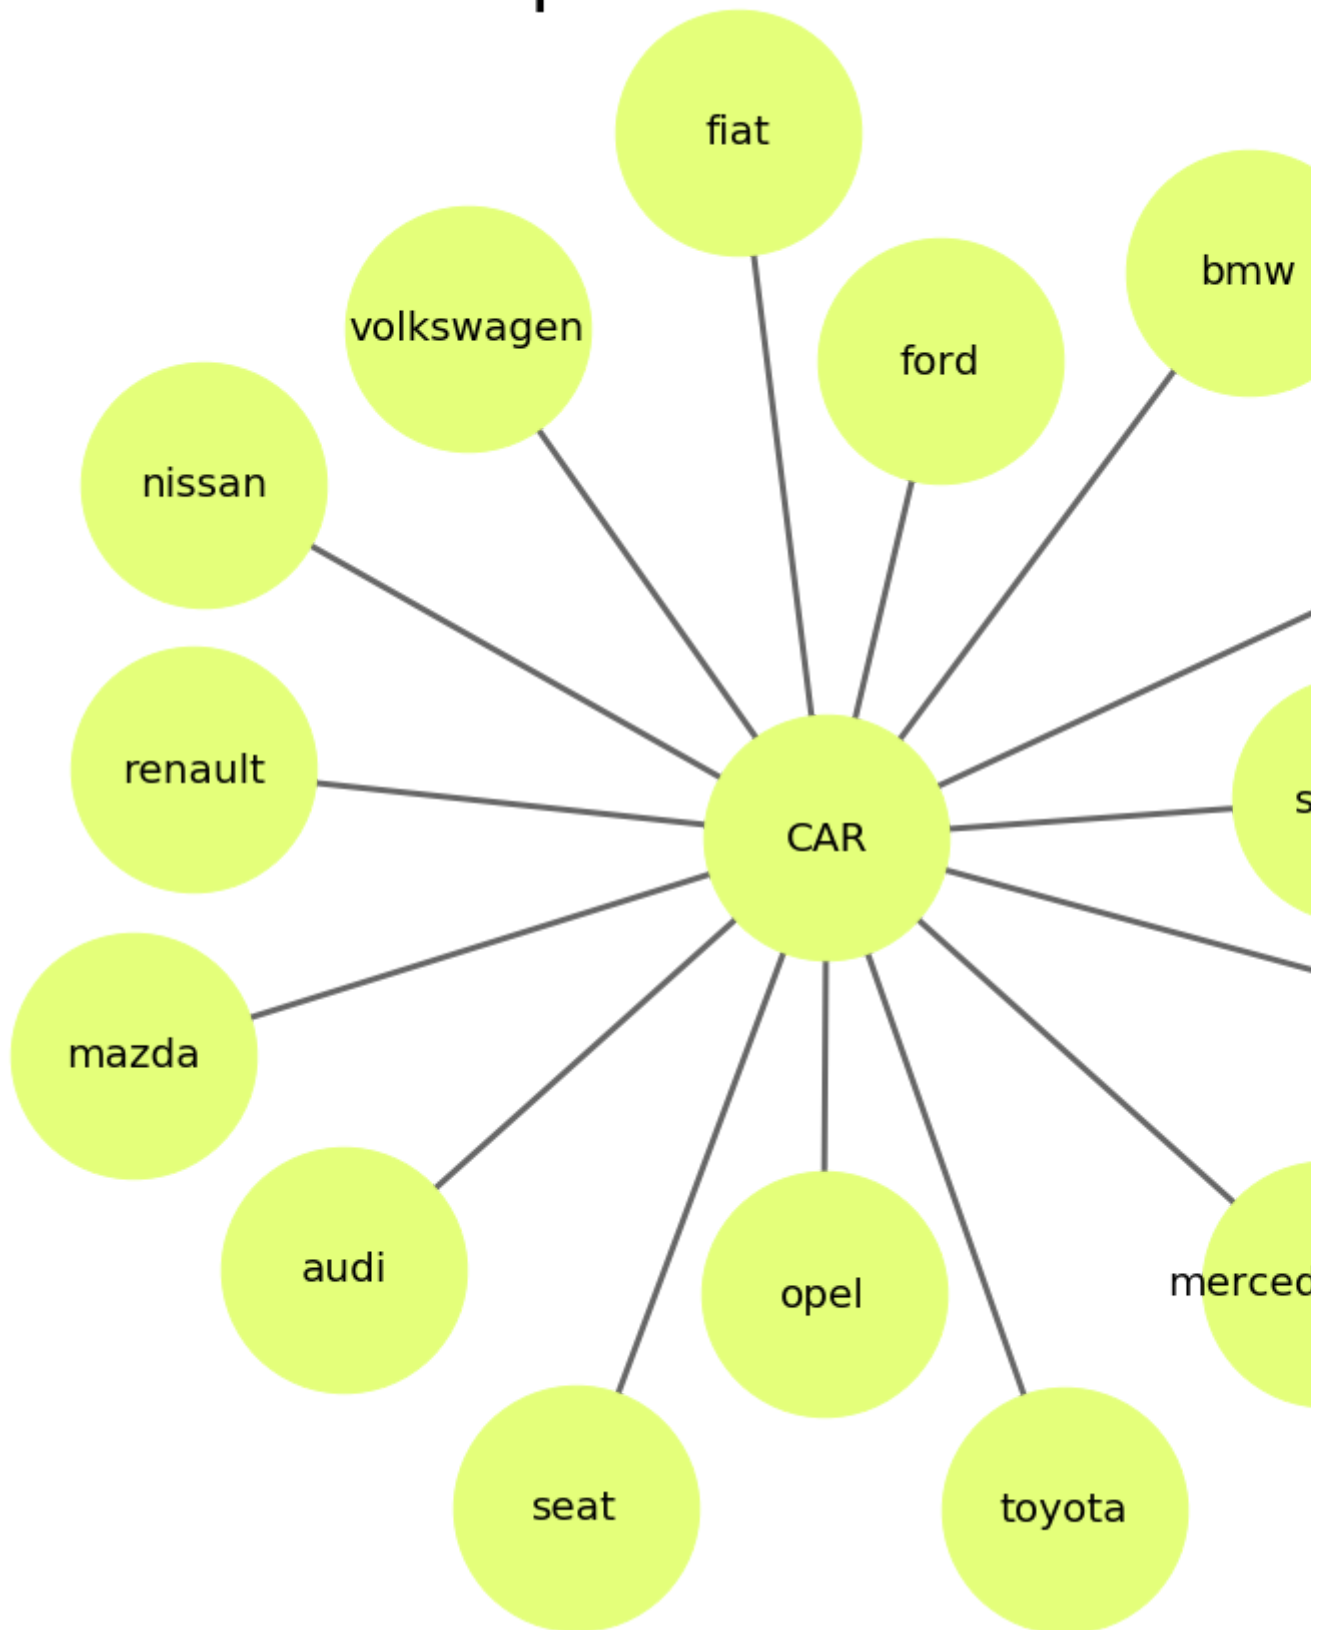
```
Car = nx.from_pandas_edgelist(Car, source = 'Car', target = 'Brand',
```

```
# Visualising the top 15 brand
import warnings
warnings.filterwarnings('ignore')
```

```
plt.rcParams['figure.figsize'] = (15, 15)
pos = nx.spring_layout(Car)
color = plt.cm.Wistia(np.linspace(0, 15, 1))
nx.draw_networkx_nodes(Car, pos, node_size = 15000, node_color = col
nx.draw_networkx_edges(Car, pos, width = 3, alpha = 0.6, edge_color
nx.draw_networkx_labels(Car, pos, font_size = 20, font_family = 'sar
plt.axis('off')
plt.grid()
plt.title('Top 15 First Choices', fontsize = 40)
plt.show()
```



Top 15 First Choices



```
# Correlation table, feature and feature correlation  
corMat = df.corr(method='pearson')
```

```
corMat
```

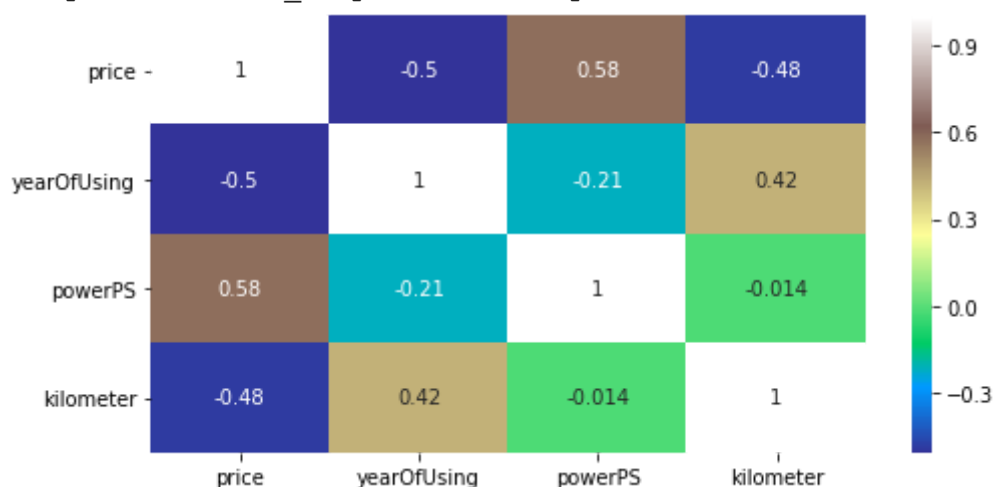


	price	yearOfUsing	powerPS	kilometer
price	1.000000	-0.504308	0.581624	-0.484205
yearOfUsing	-0.504308	1.000000	-0.211585	0.424714
powerPS	0.581624	-0.211585	1.000000	-0.013872
kilometer	-0.484205	0.424714	-0.013872	1.000000

```
# Heat map_positive and negative correlation
plt.figure(figsize=(8,4))
sns.heatmap(corMat, annot=True,cmap="terrain", )
```



```
<matplotlib.axes._subplots.AxesSubplot at 0x7f11068a3278>
```



clean data

```
# Review all column index
df.columns.values
```



```
array(['dateCrawled', 'name', 'price', 'vehicleType',
       'yearOfRegistration', 'yearOfUsing', 'gearbox', 'powerPS', 'model',
       'kilometer', 'monthOfRegistration', 'fuelType', 'brand',
       'notRepairedDamage', 'dateCreated', 'postalCode', 'lastSeen'],
      dtype=object)
```

```
# Drop not important columns
df_new = df.drop(['dateCrawled', 'name', 'yearOfRegistration', 'mont
df_new.columns.values
```



```
array(['price', 'vehicleType', 'yearOfUsing', 'gearbox', 'powerPS',
      'model', 'kilometer', 'fuelType', 'brand', 'notRepairedDamage',
      'postalCode'], dtype=object)
```

```
# Label Encoding
```

```
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
```

```
columns = ['vehicleType', 'gearbox', 'model', 'fuelType', 'brand', 'notRepairedDamage', 'postalCode']
```

```
le = LabelEncoder()
```

```
df_new[columns] = df_new[columns].apply(lambda x: le.fit_transform(x))
```

```
df_new.head()
```

```
↵
```

	price	vehicleType	yearOfUsing	gearbox	powerPS	model	kilometer	fuelType
0	2238	5	13	1	75	69	125000	6
1	12500	6	9	1	170	82	150000	1
2	13299	2	10	1	218	5	125000	6
3	17200	6	4	1	184	8	150000	1
4	8000	2	36	1	218	10	150000	6

```
df_new.info()
```

```
↵ <class 'pandas.core.frame.DataFrame'>
Int64Index: 242105 entries, 0 to 242108
Data columns (total 11 columns):
price                242105 non-null int64
vehicleType          242105 non-null int64
yearOfUsing          242105 non-null int64
gearbox              242105 non-null int64
powerPS              242105 non-null int64
model                242105 non-null int64
kilometer            242105 non-null int64
fuelType             242105 non-null int64
brand                242105 non-null int64
notRepairedDamage    242105 non-null int64
postalCode           242105 non-null int64
dtypes: int64(11)
memory usage: 32.2 MB
```

▼ Split Dataset

```
from sklearn import neighbors
```

```

from sklearn.metrics import mean_squared_error
from math import sqrt

from sklearn.model_selection import train_test_split
#feature
X= df_new.drop(['price'], axis=1)
#lable
y =df_new['price']

X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0

```

▼ KNN_Regression

```

Checnk K from 1 to 20
rmse_val = []
for K in range(20):
    K = K+1
    model = neighbors.KNeighborsRegressor(n_neighbors = K)
    model.fit(X_train, y_train)
    y_pred=model.predict(X_test)
    error = sqrt(mean_squared_error(y_test,y_pred))
    rmse_val.append(error)
    #RootMeanSquareError get avarage error
    print('RMSE value for k= ', K , 'is:', error)

Then k = which number we can get best modle
print('\nWhen k=',rmse_val.index(min(rmse_val))+1, 'we can get minnum

```



```

RMSE value for k= 1 is: 4979.083620730316
RMSE value for k= 2 is: 4639.4695260430935
RMSE value for k= 3 is: 4511.547261620518
RMSE value for k= 4 is: 4457.800532148259
RMSE value for k= 5 is: 4446.4365010197225
RMSE value for k= 6 is: 4446.592946786112
RMSE value for k= 7 is: 4444.9174576671785
RMSE value for k= 8 is: 4444.5148207438915
RMSE value for k= 9 is: 4447.485179544738
RMSE value for k= 10 is: 4455.54721489346
RMSE value for k= 11 is: 4467.10948983663
RMSE value for k= 12 is: 4479.778886598858
RMSE value for k= 13 is: 4495.478204768784
RMSE value for k= 14 is: 4506.629204576699
RMSE value for k= 15 is: 4518.263685895295
RMSE value for k= 16 is: 4532.454073668732
RMSE value for k= 17 is: 4548.850189172047
RMSE value for k= 18 is: 4561.336844808742
RMSE value for k= 19 is: 4572.244096040852
RMSE value for k= 20 is: 4582.587956341021

```

When k= 8 we can get minmum RMSE: 4444.5148207438915

▼ Linear Regression

```

import numpy as np
import pandas as pd
import matplotlib as mpl
import matplotlib.pyplot as plt
plt.figure
import seaborn as sns

from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from math import sqrt

#feature
X = df_new.drop(['price'], axis=1)
#label
y = df_new['price']

X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.

regressor=LinearRegression()
regressor.fit(X_train,y_train)
# One Hot Encoding_linear regression
df_new = pd.get_dummies(df_new[['price', 'vehicleType', 'yearOfUsing',

```

```
↳ LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
#predict price
#y=ax+b, intercept is b,slope is a
model = LinearRegression().fit(X_train, y_train)
y_pred = regressor.predict(X_test)
errorS = sqrt(mean_squared_error(y_test,y_pred))
```

```
print('RMSE:',errorS)
print('R^2:', model.score(X_test, y_test))
```

```
#To retrieve the intercept:
print('Intercept:', model.intercept_)
```

```
#For retrieving the slope :
```

```
print('Coefficient:',model.coef_)
```

```
↳ RMSE: 4560.420858721138
   R^2: 0.6342995799196617
   Intercept: 13834.521737543158
   Coefficient: [-1.35050750e+02 -2.07370718e+02 -9.41232593e+02  6.14295267e+01
                 5.34026364e+00 -7.92053137e-02 -4.81915952e+02  3.86127342e+00
                 -1.96725512e+03  5.53021008e-02]
```

```
regressor.predict(X_train)
```

```
↳ array([ 5706.19333666, -284.08916516, 31027.94579734, ...,
          15419.4043989 , 14020.1718278 , 1319.46062908])
```

```
X_test_plot= X_test['kilometer']
```

```
mpl.matplotlib_fname()
```

```
↳ '/usr/local/lib/python3.6/dist-packages/matplotlib/mpl-data/matplotlibrc'
```

```
# Visualising the Test set results
plt.figure(figsize=(6,4))
plt.scatter(y_test,y_pred, s = 20, alpha = 0.1)
plt.title('Estimate Price vs.Real Price')
plt.xlabel('Real Price')
plt.ylabel('Predicted Price')
```

```
plt.plot([min(y_test),max(y_test)],[min(y_test),max(y_test)], 'r')
#plt.plot(X_train, regressor.predict(X_train), color = 'blue')
```

```
plt.tight_layout()
#plt.show()
```



▼ Random Forest

```
#from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
import matplotlib.pyplot as plt
```

```
rfr = RandomForestRegressor(n_estimators=10)
rfr.fit(X_train, y_train)
rfr.score(X_test, y_test)
```



```
0.8872268298873156
```

```
y_pred = rfr.predict(X_test)
```

```
plt.figure(figsize=(6,4))
plt.scatter(y_test,y_pred,s=20,alpha=0.1)
plt.title('Estimated price vs. real Price')
plt.xlabel('Real Price')
plt.ylabel('Predicted Price')
```

```
plt.plot([min(y_test),max(y_test)], [min(y_test),max(y_test)], 'r')
plt.tight_layout()
```





Reference

- https://chrisalbon.com/python/data_wrangling/pandas_list_unique_values_in_column/
- <https://cmdlinetips.com/2019/10/how-to-drop-rows-based-on-a-column-value-in-pandas-dataframe/>
- <https://www.analyticsvidhya.com/blog/2018/08/k-nearest-neighbor-introduction-regression-python/>
- <https://www.youtube.com/watch?v=sA1K22Hmh1g>